

COMP 494 Final Project

Author: Hiromi Gonzalez, Camden Engstrom, Jacob Weil, Nic Harrison, Enrique Garza

Date: May 22, 2023

Bank Marketing

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be (or not) subscribed.

The dataset is ordered by date (from May 2008 to November 2010).

Table of Contents:

- [Data Importing and Pre-processing](#)
- [Data Analysis and Visualization](#)
- [Data Analytics](#)

Data Importing and Pre-processing

```
In [139... # import libraries needed
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn import tree
```

```
In [75]: # read in file
df = pd.read_csv('bank_marketing.csv', sep=";")
```

```
In [76]: df.default.value_counts()
```

```
Out[76]: no      43113
yes       792
Name: default, dtype: int64
```

```
In [77]: #check number of rows and columns
df.shape #45,211 rows and 17 columns
```

```
Out[77]: (45211, 17)
```

```
In [78]: #count the number of categorical variables
cat_count = 0
for dtype in df.dtypes:
    if dtype == 'object':
        cat_count = cat_count + 1
```

```
In [79]: print('# of categorical variables:',cat_count)
print('# of continuous variables:',df.shape[1] - cat_count)
```

```
# of categorical variables: 10
# of continuous variables: 7
```

```
In [80]: #check the column names
df.columns
```

```
Out[80]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
               'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays',
               'previous', 'poutcome', 'deposit'],
              dtype='object')
```

```
In [83]: # checking data types of each columns
column_data_types = df.dtypes
print(column_data_types)
```

```
age          float64
job           object
marital       object
education     object
default       object
balance       int64
housing       object
loan          object
contact       object
day           int64
month         object
duration      int64
campaign      int64
pdays        int64
previous      int64
poutcome     object
deposit       object
dtype: object
```

Handling missing data

```
In [81]: #missing data (as percentages)
cleaned_df = df
total = df.isnull().sum().sort_values(ascending=False)
percent = (df.isnull().sum()/df.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(20)
```

Out [81]:

	Total	Percent
contact	1383	0.030590
age	1339	0.029617
default	1306	0.028887
marital	0	0.000000
education	0	0.000000
balance	0	0.000000
housing	0	0.000000
loan	0	0.000000
job	0	0.000000
day	0	0.000000
month	0	0.000000
duration	0	0.000000
campaign	0	0.000000
pdays	0	0.000000
previous	0	0.000000
poutcome	0	0.000000
deposit	0	0.000000

In [82]: *# Find the number of nulls for each column*

```
null_counts = df.isnull().sum()
print(null_counts)
```

```
age          1339
job           0
marital       0
education     0
default      1306
balance       0
housing       0
loan          0
contact      1383
day           0
month         0
duration      0
campaign      0
pdays        0
previous      0
poutcome      0
deposit       0
dtype: int64
```

In [84]: *# inserting the average age of column into the null values of age*

```
average_age = df['age'].mean()
cleaned_df['age'].fillna(average_age, inplace=True)
```

```
In [85]: # assigning "unknown" to all contact that have null values  
cleaned_df['contact'].fillna("unknown", inplace=True)
```

```
In [86]: # assigning "no" to all values in default column that are NaN  
cleaned_df['default'].fillna("no", inplace=True)
```

```
In [87]: null_counts = cleaned_df.isnull().sum()  
print(null_counts)
```

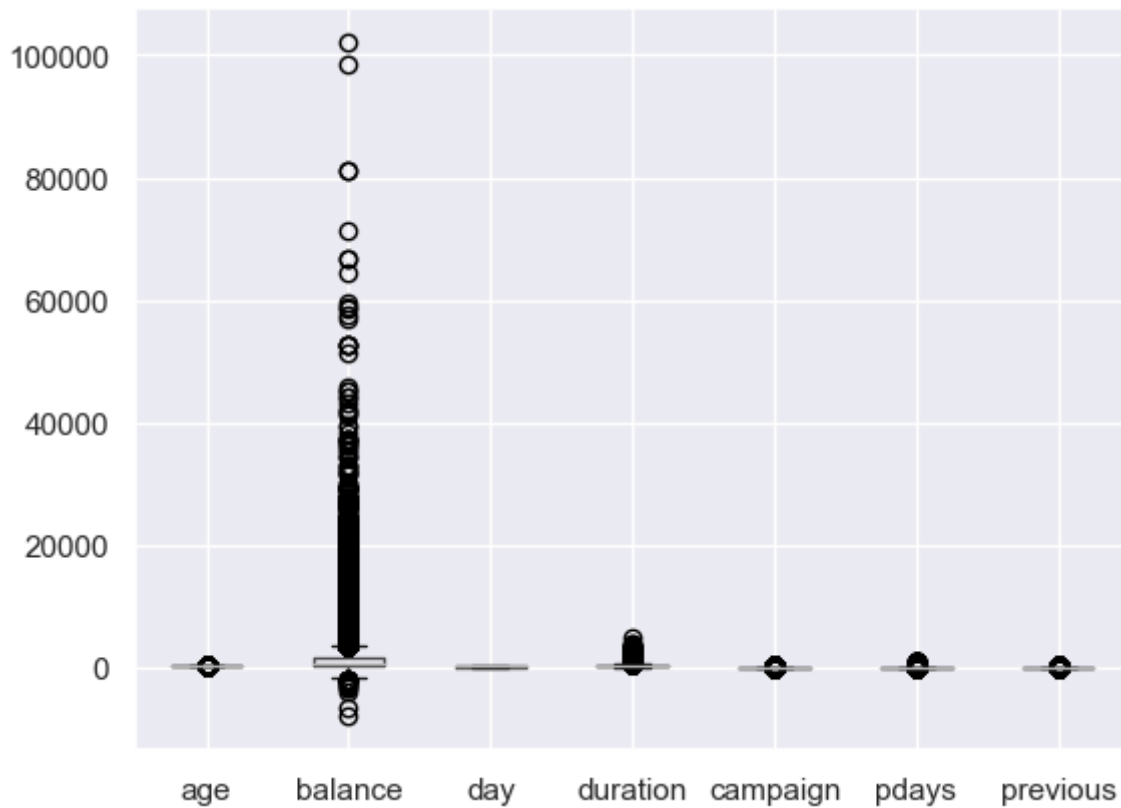
```
age          0  
job          0  
marital      0  
education    0  
default      0  
balance      0  
housing      0  
loan         0  
contact      0  
day          0  
month        0  
duration     0  
campaign     0  
pdays       0  
previous     0  
poutcome     0  
deposit      0  
dtype: int64
```

```
In [88]: cleaned_df.shape
```

```
Out[88]: (45211, 17)
```

Handling Outliers

```
In [91]: # create boxplot showing all the outliers  
cleaned_df.boxplot()  
plt.show()  
cleaned_df.describe()
```



Out[91]:

	age	balance	day	duration	campaign	pdays
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.924781	1362.272058	15.806419	258.163080	2.763841	40.197828
std	10.452521	3044.765829	8.322476	257.527812	3.098021	100.128746
min	18.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000	-1.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000	-1.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000	871.000000

```
In [92]: # do not remove outliers
cleaned_without_outliers_df = cleaned_df.copy()
```

```
In [94]: summary_stats = df.describe()

print(summary_stats)
```

	age	balance	day	duration	campaign
\					
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.924781	1362.272058	15.806419	258.163080	2.763841
std	10.452521	3044.765829	8.322476	257.527812	3.098021
min	18.000000	-8019.000000	1.000000	0.000000	1.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000

	pdays	previous
count	45211.000000	45211.000000
mean	40.197828	0.580323
std	100.128746	2.303441
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	-1.000000	0.000000
max	871.000000	275.000000

Making necessary columns to be binary values

```
In [95]: # changing deposit to binary variable (no==0, yes==1)
cleaned_without_outliers_df['deposit'] = cleaned_df['deposit'].map({'no': 0, 'yes': 1})
```

```
In [96]: # changing other categorical variables to be binary (no==0, yes==1)
cleaned_without_outliers_df['default'] = cleaned_df['default'].map({'no': 0, 'yes': 1})
cleaned_without_outliers_df['housing'] = cleaned_df['housing'].map({'no': 0, 'yes': 1})
cleaned_without_outliers_df['loan'] = cleaned_df['loan'].map({'no': 0, 'yes': 1})
```

```
In [97]: cleaned_without_outliers_df['housing'].head(5)
```

```
Out[97]: 0    1
1    1
2    1
3    1
4    0
Name: housing, dtype: int64
```

Finished data cleaning of the Bank_Marketing.csv file - moving onto analysis now

Data Analysis and Visualization

```
In [100]: # installing necessary libraries/modules
```

```
In [101]: !pip3 install xgboost
```

```
Requirement already satisfied: xgboost in /Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages (1.7.5)
Requirement already satisfied: numpy in /Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages (from xgboost) (1.21.5)
Requirement already satisfied: scipy in /Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages (from xgboost) (1.9.1)
```

```
In [24]: cleaned_without_outliers_df.head()
```

Out[24]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	58.0	management	married	tertiary	0	2143	1	0	unknown	5	may
1	44.0	technician	single	secondary	0	29	1	0	unknown	5	may
2	33.0	entrepreneur	married	secondary	0	2	1	1	unknown	5	may
3	47.0	blue-collar	married	unknown	0	1506	1	0	unknown	5	may
4	33.0	unknown	single	unknown	0	1	0	0	unknown	5	may

Encoding

```
In [109... # label encoding
cleaned_without_outliers_df.job.value_counts()
```

```
Out[109]: blue-collar      9732
management    9458
technician     7597
admin.         5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student        938
unknown        288
Name: job, dtype: int64
```

```
In [110... # one hot encoding
cleaned_without_outliers_df.marital.value_counts()
```

```
Out[110]: married      27214
single      12790
divorced     5207
Name: marital, dtype: int64
```

```
In [111... # one hot encode
cleaned_without_outliers_df.education.value_counts()
```

```
Out[111]: secondary    23202
tertiary    13301
primary     6851
unknown     1857
Name: education, dtype: int64
```

```
In [112... # one hot encode
cleaned_without_outliers_df.contact.value_counts()
```

```
Out[112]: cellular    28410
unknown    13992
telephone   2809
Name: contact, dtype: int64
```

```
In [113... # manually label them and sure labels are in the proper order
cleaned_without_outliers_df.month.value_counts()
```

```
Out[113]: may      13766
          jul      6895
          aug      6247
          jun      5341
          nov      3970
          apr      2932
          feb      2649
          jan      1403
          oct       738
          sep       579
          mar       477
          dec       214
          Name: month, dtype: int64
```

```
In [114... # one hot encode
cleaned_without_outliers_df.poutcome.value_counts()
```

```
Out[114]: unknown    36959
          failure    4901
          other      1840
          success    1511
          Name: poutcome, dtype: int64
```

```
In [115... # try to predict deposit column
          # based on the info we have we are trying to predict if they will make a deposit
```

```
In [116... new_df = cleaned_without_outliers_df.copy()
```

```
In [117... # look up Classifier rather than regression, the different
new_df.head()
```

```
Out[117]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	58.0	management	married	tertiary	0	2143	1	0	unknown	5	may
1	44.0	technician	single	secondary	0	29	1	0	unknown	5	may
2	33.0	entrepreneur	married	secondary	0	2	1	1	unknown	5	may
3	47.0	blue-collar	married	unknown	0	1506	1	0	unknown	5	may
4	33.0	unknown	single	unknown	0	1	0	0	unknown	5	may

```
In [118... month_labels = {
    'jan': 1,
    'feb': 2,
    'mar': 3,
    'apr': 4,
    'may': 5,
    'jun': 6,
    'jul': 7,
    'aug': 8,
    'sep': 9,
    'oct': 10,
    'nov': 11,
    'dec': 12
}

new_df['month'] = new_df['month'].map(month_labels)
```


In [119... `new_df.head()`

Out[119]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	58.0	management	married	tertiary	0	2143	1	0	unknown	5	
1	44.0	technician	single	secondary	0	29	1	0	unknown	5	
2	33.0	entrepreneur	married	secondary	0	2	1	1	unknown	5	
3	47.0	blue-collar	married	unknown	0	1506	1	0	unknown	5	
4	33.0	unknown	single	unknown	0	1	0	0	unknown	5	

In [120... `encoder = LabelEncoder()`
`new_df['job'] = encoder.fit_transform(new_df['job'])`

In [121... `new_df.head()`

Out[121]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration
0	58.0	4	married	tertiary	0	2143	1	0	unknown	5	5	
1	44.0	9	single	secondary	0	29	1	0	unknown	5	5	
2	33.0	2	married	secondary	0	2	1	1	unknown	5	5	
3	47.0	1	married	unknown	0	1506	1	0	unknown	5	5	
4	33.0	11	single	unknown	0	1	0	0	unknown	5	5	

In [122... `new_df = pd.get_dummies(new_df, columns=['marital', 'education', 'contact', 'pou`

In [123... `new_df.head()`


Out[123]:

	age	job	default	balance	housing	loan	day	month	duration	campaign	...	education_
0	58.0	4	0	2143	1	0	5	5	261	1	...	
1	44.0	9	0	29	1	0	5	5	151	1	...	
2	33.0	2	0	2	1	1	5	5	76	1	...	
3	47.0	1	0	1506	1	0	5	5	92	1	...	
4	33.0	11	0	1	0	0	5	5	198	1	...	

5 rows x 27 columns


Target Variable Scatterplots

In [124... `sns.set()`
`cols = ['age', 'job', 'default', 'housing', 'loan', 'day', 'month', 'duration', 'campaign', 'education_tertiary', 'education_unknown', 'contact_cellular', 'contact_outcome_failure', 'outcome_other', 'outcome_success', 'outcome_unknown', 'deposit']`
`sns.pairplot(new_df[['age', 'job', 'deposit', 'loan', 'duration', 'campaign']],`
`plt.show())`
the more people that were in a campaign were more likely to open a deposit



Correlation Matrix

```
In [126... #Correlation map to see how features are correlated
corrmatrix = new_df.corr()
f, ax = plt.subplots(figsize=(15, 12))
sns.heatmap(corrmatrix, vmax=.8, square=True)
plt.show()
```



Logistic Regression

```
In [133... # Split the data into features (X) and the target variable (y)
X = new_df[['age', 'job', 'default', 'housing', 'loan', 'duration', 'campaign', 'educ
            'education_tertiary', 'education_unknown', 'contact_cellular', 'contact_
            'poutcome_failure', 'poutcome_other', 'poutcome_success', 'poutcome_unkr
y = new_df[['deposit']]
```

```
In [134... X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random
```

```
In [135... model = LogisticRegression()
```

```
In [136... model.fit(X_train, y_train)
```

```
/Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/
validation.py:993: DataConversionWarning: A column-vector y was passed when a
1d array was expected. Please change the shape of y to (n_samples, ), for exam
ple using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
/Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear
_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=
1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Out[136]: LogisticRegression()

In [140... *# Make predictions on the test set*

```
y_pred = model.predict(X_test)
```

```
# Evaluate the model
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print(classification_report(y_test, y_pred))
```

```
[[7758  194]
```

```
 [ 747  344]]
```

	precision	recall	f1-score	support
0	0.91	0.98	0.94	7952
1	0.64	0.32	0.42	1091
accuracy			0.90	9043
macro avg	0.78	0.65	0.68	9043
weighted avg	0.88	0.90	0.88	9043

These metrics indicate that the model performs better for predicting when there won't be a deposit (0), as it has higher precision, recall, and F1-score compared to predicting when there will be a deposit (1).

In [148... *# Get the coefficients and their corresponding feature names*

```
coefficients = model.coef_[0]
```

```
feature_names = X.columns
```

```
# Create a bar plot
```

```
plt.bar(feature_names, coefficients)
```


```
plt.xlabel('Features')
```

```
plt.ylabel('Coefficients')
```

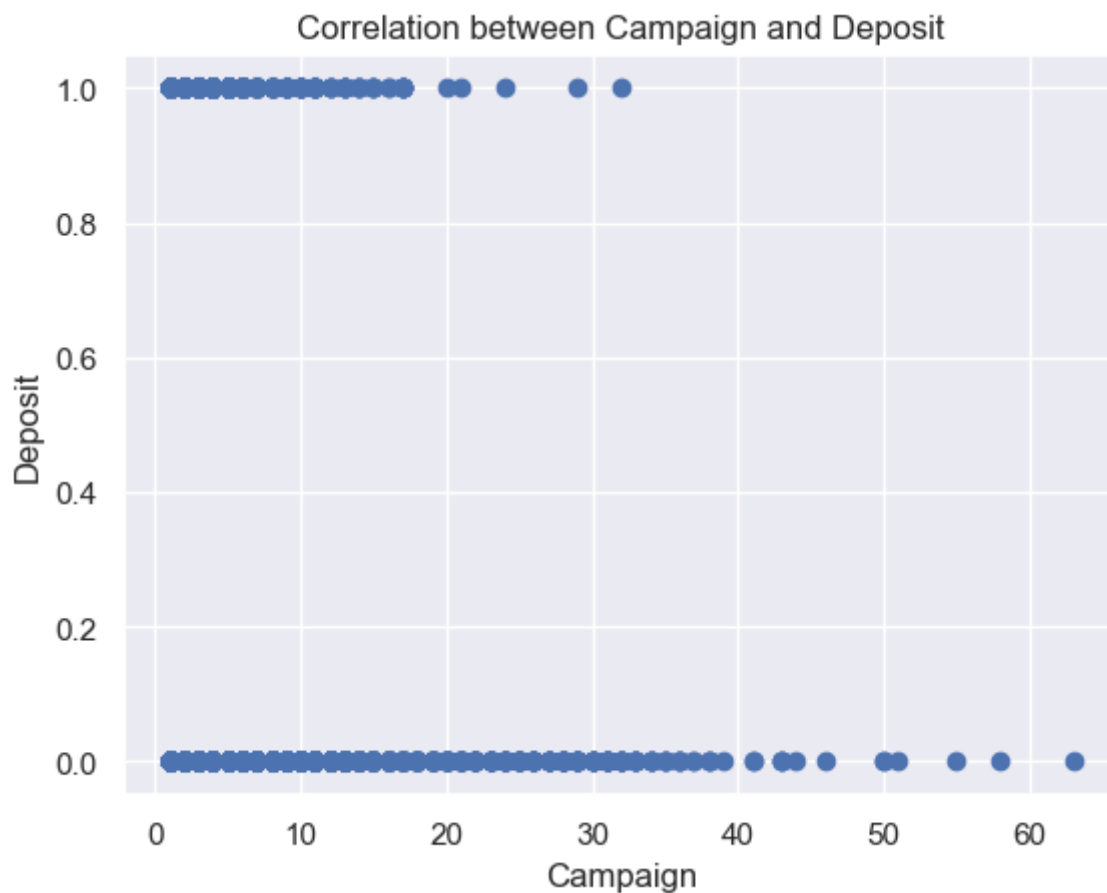
```
plt.title('Logistic Regression Coefficients')
```

```
plt.xticks(rotation=90) # Rotate x-axis labels if needed for better readability
```

```
plt.show()
```



```
In [46]: deposit = new_df['deposit']
campaign = new_df['campaign']
plt.scatter(campaign, deposit)
plt.xlabel('Campaign')
plt.ylabel('Deposit')
plt.title('Correlation between Campaign and Deposit')
plt.show()
```



Data Analytics

```
In [149... from sklearn.linear_model import Lasso, LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_val_score
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
import xgboost as xgb
from sklearn.metrics import accuracy_score
```

```
import xgboost as xgb from sklearn.metrics import accuracy_score
```

Init classifier

```
xgb_cl = xgb.XGBClassifier()
```

Fit

```
xgb_cl.fit(X_train, y_train)
```

Predict

```
preds = xgb_cl.predict(X_test)
```

```
In [150... X = new_df[['age', 'job', 'default', 'housing', 'loan', 'duration', 'campaign', 'educ
            'education_tertiary', 'education_unknown', 'contact_cellular', 'contact_
            'poutcome_failure', 'poutcome_other', 'poutcome_success', 'poutcome_unkr
```

```
y = new_df[['deposit']]

# create classification model and fit the data
xgb_cl = xgb.XGBClassifier()
```

```
In [151...] xgb_cl.fit(X,y)
```

```
Out[151]: XGBClassifier(base_score=None, booster=None, callbacks=None,
      colsample_bylevel=None, colsample_bynode=None,
      colsample_bytree=None, early_stopping_rounds=None,
      enable_categorical=False, eval_metric=None, feature_types=None,
      gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
      interaction_constraints=None, learning_rate=None, max_bin=None,
      max_cat_threshold=None, max_cat_to_onehot=None,
      max_delta_step=None, max_depth=None, max_leaves=None,
      min_child_weight=None, missing=nan, monotone_constraints=None,
      n_estimators=100, n_jobs=None, num_parallel_tree=None,
      predictor=None, random_state=None, ...)
```

```
In [152...] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, rand
```

```
In [153...] model = xgb.XGBClassifier()
```

```
In [154...] model.fit(X_train, y_train)
```

```
Out[154]: XGBClassifier(base_score=None, booster=None, callbacks=None,
      colsample_bylevel=None, colsample_bynode=None,
      colsample_bytree=None, early_stopping_rounds=None,
      enable_categorical=False, eval_metric=None, feature_types=None,
      gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
      interaction_constraints=None, learning_rate=None, max_bin=None,
      max_cat_threshold=None, max_cat_to_onehot=None,
      max_delta_step=None, max_depth=None, max_leaves=None,
      min_child_weight=None, missing=nan, monotone_constraints=None,
      n_estimators=100, n_jobs=None, num_parallel_tree=None,
      predictor=None, random_state=None, ...)
```

```
In [155...] # make the prediction on test
y_test_prediction = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_test_prediction))
print(f"RMSE: {rmse}")
```

RMSE: 0.3248014219932712

```
In [156...] # now fine the RMSE of the train set
y_train_predict = model.predict(X_train)

# calculate the RMSE on the train set now
rmse_train = np.sqrt(mean_squared_error(y_train, y_train_predict))
print(f"RMSE train: {rmse_train}")
```

RMSE train: 0.26505749273045415

-from sklearn.model_selection import KFold, cross_val_score

-from sklearn.tree import DecisionTreeClassifier

-from sklearn.neighbors import KNeighborsClassifier

```
In [157... clf = RandomForestClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy_random_forest_classifier = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy_random_forest_classifier}")
```

/var/folders/54/937x7kcn6m79plvcfs7pl6zw0000gn/T/ipykernel_36408/4237262992.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
    clf.fit(X_train, y_train)
Accuracy: 0.8904124737365918
```

```
In [158... # now we will run a decision tree classifier test
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy_decision_tree_classifier = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy_decision_tree_classifier}")
```

```
Accuracy: 0.8500497622470419
```

```
In [159... clf = KNeighborsClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy_kneighbors_classifier = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy_kneighbors_classifier}")
```

/Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:198: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self._fit(X, y)

/Users/hiromigonzalez/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

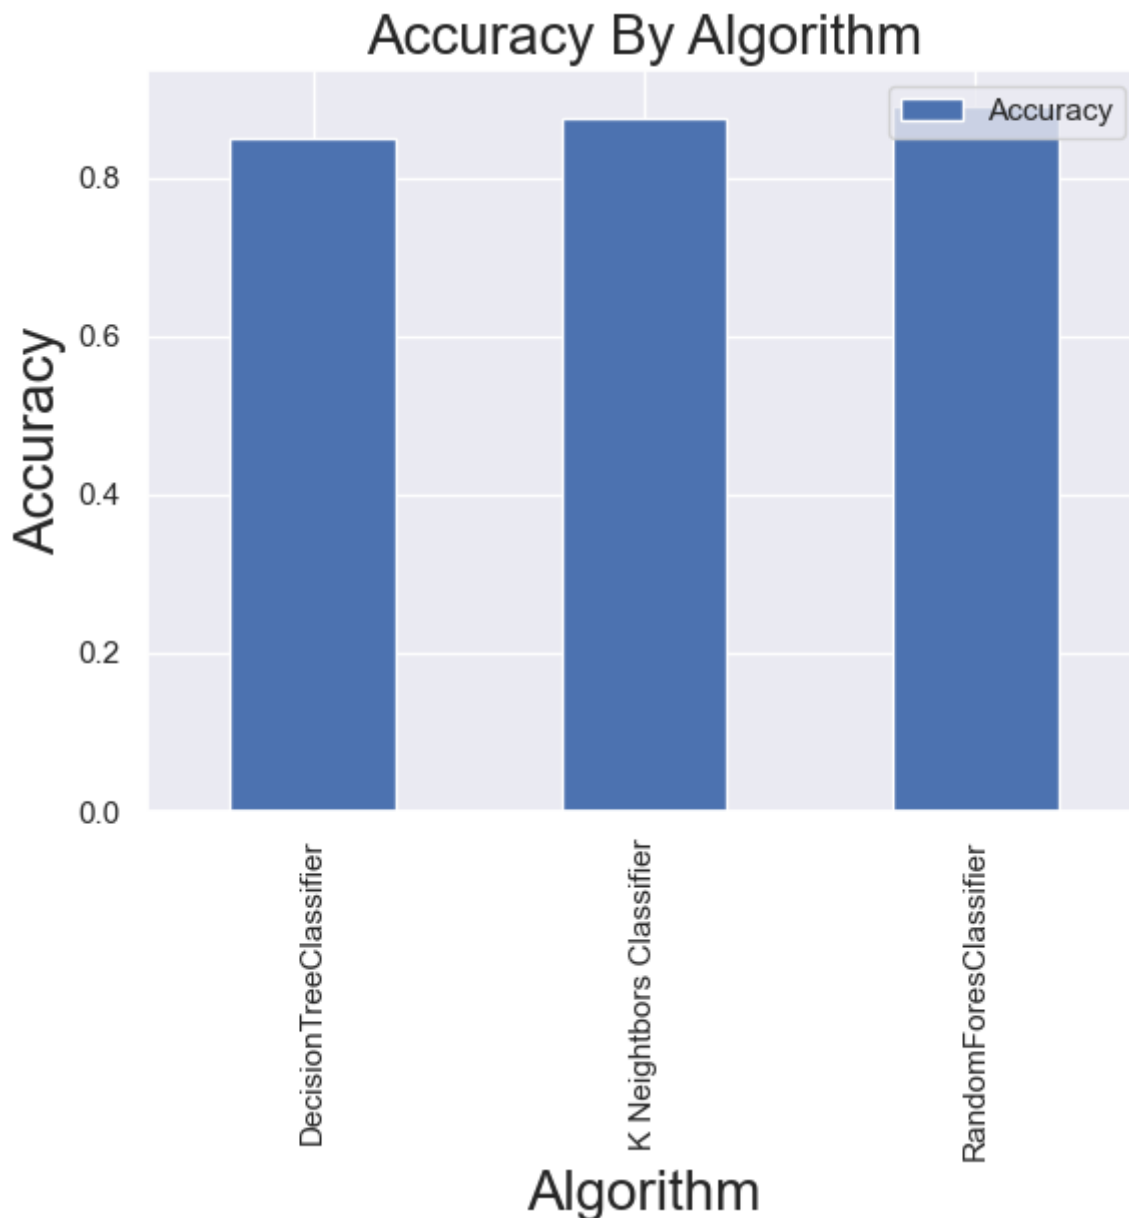
```
Accuracy: 0.8752626340816101
```

```
In [160... # plot accuracy for each algorithm
data = {'RandomForesClassifier': [accuracy_random_forest_classifier], 'DecisionT
      'K Neighbors Classifier': [accuracy_kneighbors_classifier]}
data_df = pd.DataFrame(data=data).T.reset_index().sort_values(by = [0], ascending=True)
data_df.columns = ['Algorithm', 'Accuracy']
data_df.head()
```

```
Out[160]:
```

	Algorithm	Accuracy
1	DecisionTreeClassifier	0.850050
2	K Neighbors Classifier	0.875263
0	RandomForesClassifier	0.890412


```
In [161... data_df.plot(kind='bar',x = 'Algorithm', y = 'Accuracy')
plt.xlabel("Algorithm",fontsize=20)
plt.ylabel("Accuracy",fontsize=20)
plt.title("Accuracy By Algorithm",fontsize=20)
plt.show()
```



Variable Importance Plot

```
In [162... clf = RandomForestClassifier()
clf.fit(X_train, y_train)
importances = clf.feature_importances_
# Create a DataFrame with feature names and importances
importance_df = pd.DataFrame({'Feature': X_train.columns, 'Importance': importances})
# Sort the DataFrame by importance values in descending order
importance_df = importance_df.sort_values(by='Importance', ascending=False)
# Plot the variable importance
plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'], importance_df['Importance'])
plt.xlabel('Importance')
```

```
plt.ylabel('Feature')  
plt.title('Variable Importance')  
plt.show()
```

```
/var/folders/54/937x7kcn6m79plvcfs7pl6zw0000gn/T/ipykernel_36408/1652506768.p  
y:2: DataConversionWarning: A column-vector y was passed when a 1d array was e  
xpected. Please change the shape of y to (n_samples,), for example using ravel  
( ).  
clf.fit(X_train, y_train)
```

